A blue parallelogram and a light green parallelogram are positioned in the top-left corner of the slide. The background is a dark navy blue with several diagonal bands of a slightly lighter shade of navy blue running from the bottom-left towards the top-right.

# 10 things

I wish I knew about

# Code efficiency

as a junior dev



# This is a talk about code efficiency

- > I am not an expert
- > But I got this far without needing to be one
- > Improving efficiency is 50% technical knowledge, and 50% problem solving approaches



# What is code efficiency?

- > Low latency
- > Low resource usage
- > High throughput
- > absolutely nothing to do with how many lines of code you write

1 >>

Don't worry about it  
unless you have to





1 >>

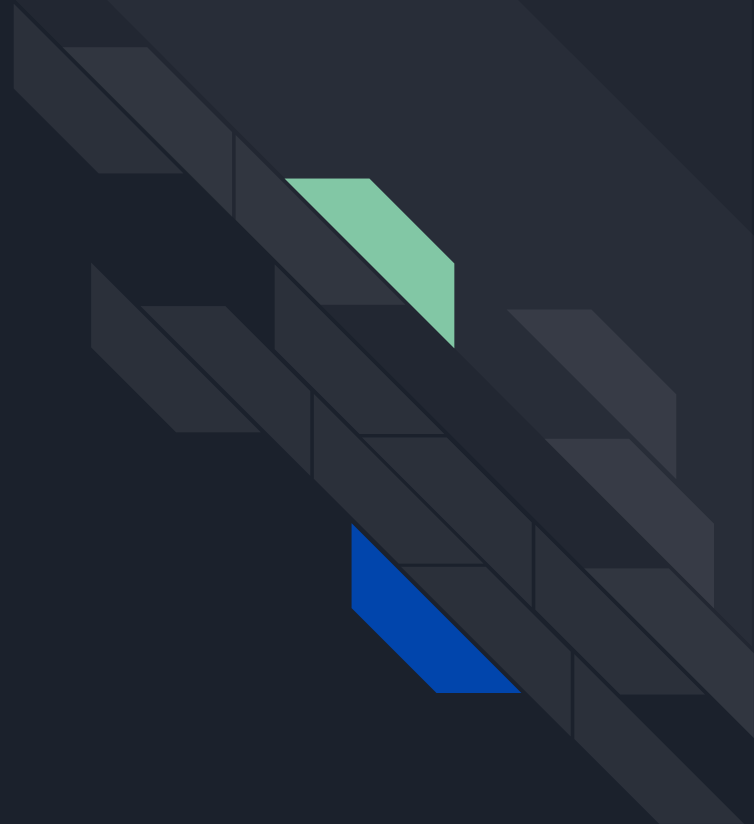
# Don't worry about it

- > You have my full permission not to use anything you learn in this talk, for, like, months
- > You've written tonnes of code, and almost all of it has been good enough
- > Make it work. Make it right. Make it fast
- > If it's a problem, you'll know
- > If you think it might be a problem in the future:
  - >> have a way to measure lag
  - >> load test your system
  - >> follow good, OOP practices

2 >>

## Find the bottleneck

20% of the code takes 80% of the processing power





2 >>

# Find your bottleneck

- > It doesn't matter how wide the rest of the bottle is, the neck is what limits the flow
- > Computer are really fast
- > If they aren't, it's probably a specific problem.
- > Don't waste time trying to make the rest of the bottle wider, fix the neck



2 >>

# Find your bottleneck

> How to find a bottleneck

>> be an expert

>> metrics, logs, and jstack

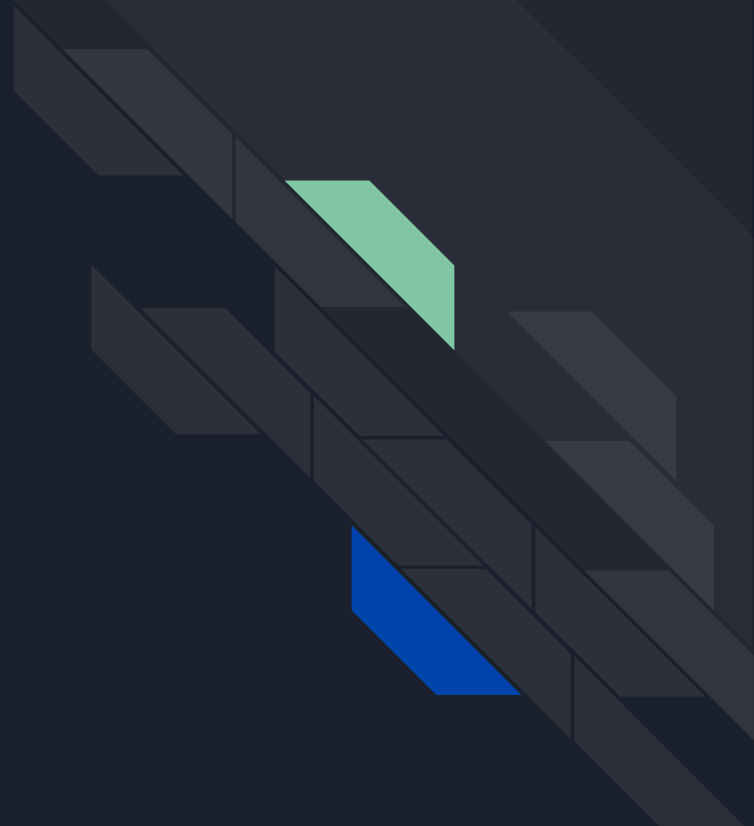
>> code profilers

>> it might not be your code. Networks and shared resources can be the slowest part



3 >>

Learn how to  
estimate efficiency  
It's not hard





3 >>

# Learn how to estimate efficiency

(Number of times you need to do X) \* (Amount of work it takes to do X)

- > Heavy processes that run occasionally are fine
- > Cheap processes you run a lot are fine
- > Heavy processes you run all the time are not fine
- > How to calculate how much work it takes to do X?



3 >>


# Learn how to estimate efficiency

Efficiency of X = number of steps it takes to do X

- > Steps, not seconds
- > The number of steps isn't always obvious
- > The number of steps isn't always constant
- > Understand what influences the number of steps, and how

3 >>

Learn how to estimate efficiency





3 >>

# Learn how to estimate efficiency

> Pop quiz! How efficient are these?

Given that `ArrayList x = [5, 11, 9243, 6...]`

`x.get(50)`

`sumList(x)` //add all the elements in x together

`x.contains(22)`

`x.containsAll(y)`

`x.sort()`

What if x was sorted?



3 >>

## Learn how to estimate efficiency

- > What about space?
- > Easy!
- > get a small sample
- > get a volume estimate
- > multiply them together

4 >>

Learn the little hacks for  
your primary coding  
language

Because it's easy to do





4 >>

# Learn the little hacks for your primary coding language

```
list.contains(x)
```

Vs

```
set.contains(x)
```





4 >>

## Learn the little hacks for your primary coding language

```
logger.debug("Processing resource: {}", resource)
```

vs

```
logger.debug("processing resource: " + resource.toString())
```



4 >>

## Learn the little hacks for your primary coding language

```
logger.debug("Processing resource: {}", resource)
```

vs

```
logger.debug("processing resource: " + resource.toString())
```



4 >>

# Learn the little hacks for your primary coding language

```
if (i == 0 || expensiveFunction())
```

vs

```
if (expensiveFunction() || i == 0)
```



4 >>

# Learn the little hacks for your primary coding language

```
Integer i = 1
```

vs

```
int i = 1
```



4 >>


# Learn the little hacks for your primary coding language

- > use caches where appropriate
- > write lazy code
- > avoid nested lists
- > avoid regexes
- > Use string builders for concatenation
- > Define collection size at creation
- > check hubot comment explanations

5 >>

Learn the deep stuff,  
too

Because there's no substitute for  
actual understanding





5 >>

## Learn the deep stuff, too

- > This isn't a lecture on Java memory management
- > This is me telling you to go to lecture on Java memory management
- > It takes a lot of time and effort to learn, but it's worth it
- > So much of efficiency is down to implementation details
- > If you don't have a basic understanding of how things work under the hood, it can be hard to even follow conversations about efficiency



5 >>

## Learn the deep stuff, too

>> What is the JVM and how does it work?

>> How is memory allocated? How much memory does your app have, and how is being assigned? How can you tune this?

>> What is garbage collection? When does it happen, and what does it do? Why is too much of it bad?

>> what are memory leaks and how to find them? What about GC thrashing? Fragmented memory?

>> what are threads? How are they managed? What states can they be in, and what do these states mean? How many threads is best?


>> what are the java primitive types, and what are their sizes? How are they different to objects? How are common collections implemented?



6 >>

Multi-threading and  
multiple instances  
provide infinite\*  
scaling

But it comes with a cost





6 >>

## Multi-threading and multiple instances provide infinite\* scaling

- > Multi-threading allows your app to execute many pieces of code at the same time
- > It's limited by the number of processors on your machine. Our live servers can ~80 processors
- > Multiple instances allow horizontal scaling. Want more power? Just deploy another instance on a new machine
- > Combine these two, and you can process truly huge amounts of data
- > But there's a cost



6 >>

## Multi-threading and multiple instances provide infinite\* scaling

- > concurrency is hard

- >> Non-linear code is harder to understand

- >> Non-deterministic bugs are harder to detect and diagnose

- >> It's not always obvious when you've made a serious mistake

- > Multiple instances of an app are hard (if they need to coordinate with each other)

- > And multiple instances of an app have another cost...



6 >>

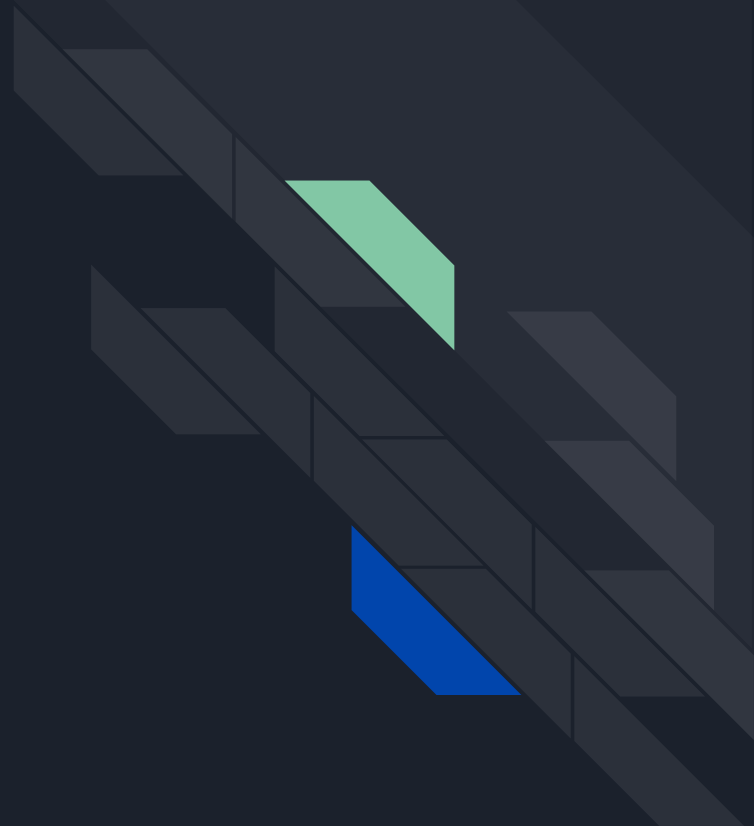
## Multi-threading and multiple instances provide infinite\* scaling

- > Hardware
- > It isn't free
- > You aren't increasing efficiency, you're increasing throughput by throwing money at the problem
- > Horizontal scaling is amazing and necessary, but it can eat all your profit if you let it
- > example: twitter firehose

7 >>

Consider redesigning  
the algorithm

It's easier than micro-optimising  
code





7 >>


# Consider redesigning the algorithm

- > Rewriting is sometimes better than improving what you have
- > Potential for massive performance gains
- > Examples
  - >> prematched query ids on tweets
  - >> mnemosyne
  - >> timestamping database entries when polling for differences
- > Requires you to really understand your problem space AND have a good idea

8 >>

If all else fails, find an  
expert

It's going to get technical





8 >>

## If all else fails, find an expert

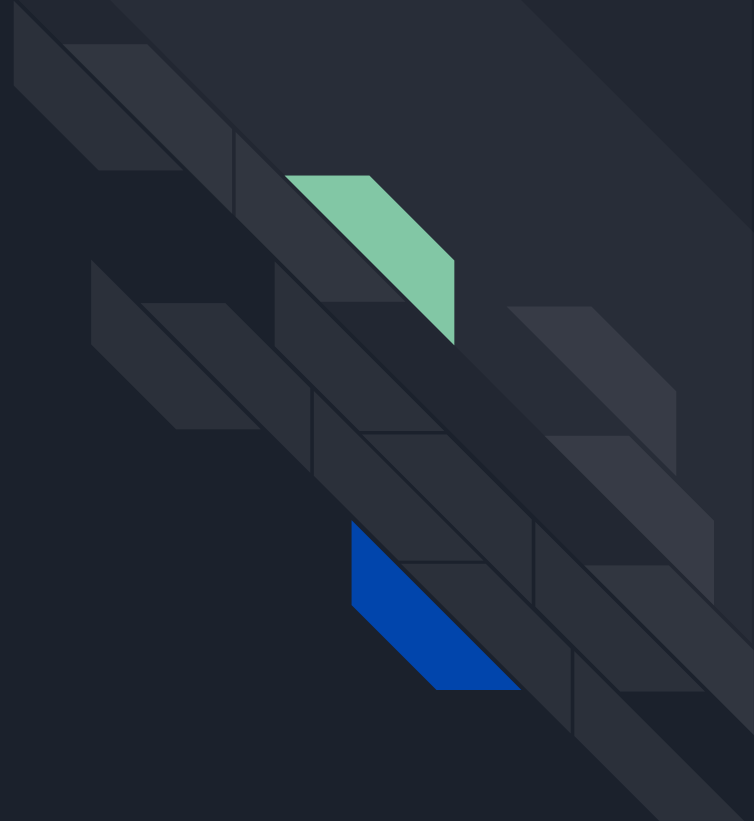
- > Micro-optimising code is a lot of work for not much gain, and it makes code less readable
- > But if you really understand what your code is doing, and really, really need it to be faster, there will be a way
- > Sometimes it's as simple as changing an int to a short
- > Sometimes it's as crazy as changing WHERE IN(...) to WHERE NOT NOT IN(...)
- > In time, you can become the expert you need
- > but you'll never be an expert in everything



9 >>

Measure your  
“efficiency  
enhancement”

Don't just assume it worked





9 >>


## Measure your “efficiency enhancement”

- > Efficiency is really complicated. Things that “should work” might not
  - >> multithreading can slow down your processes
  - >> exponential complexity can be faster than constant
  - >> improvements outside your bottleneck might not make any difference
- > Run timed comparison between your branch and the master branch

10 >>

## Be mindful of external systems

Your apps efficiency isn't the only  
one that matters





10 >>

## Be mindful of external systems

- > Just because your app can keep up with the load doesn't mean our whole system can
- > External systems might be servicing 100s of requests just like yours, and the impact adds up
- > It can be easy to overlook or not notice your impact on other systems
  - >> Databases
  - >> Networks
  - >> Services
  - >> Downstream processes



# Recap

- > If you don't have performance problems, don't worry about it
- > If you do have performance problems, focus on finding and fixing the bottle neck
- > Picking up some small efficiency hacks is easy, but has limited impact
- > Learning how stuff works under the hood is pretty important
- > Redesigning your algorithm can be better than trying to improve what you have
- > Multithreading and horizontal scaling allow massive scaling
- > hardcore black magic is also an option, but I won't recommend it